

A Novel Hybrid EEMD-TGNet Model for Accurate and Efficient Solar Energy Forecasting on Edge Devices

Anh-Phuc Pham

Department of Computer Science

University of Engineering and Technology, VNU Hanoi

Hanoi, Vietnam

anh.pham@university.edu

Abstract—The integration of solar power into smart grids demands accurate and efficient energy forecasting. However, the inherent intermittency and non-stationarity of solar radiation pose significant challenges. This paper proposes a novel hybrid deep learning framework, the Ensemble Empirical Mode Decomposition-based Temporal Convolutional Gated Recurrent Unit Network (EEMD-TGNet), to address these challenges. Our model first uses EEMD to decompose the raw solar energy time series into a set of simpler, more stationary Intrinsic Mode Functions (IMFs). Each IMF is then independently forecasted using a hybrid Temporal Convolutional Network (TCN) and Gated Recurrent Unit (GRU) model, which captures both long-range dependencies and sequential patterns. The final forecast is an aggregation of the predictions for each IMF. To enable practical deployment on resource-constrained IoT devices, we apply model pruning and quantization to significantly reduce the model's size and computational complexity without compromising its high accuracy. Experimental validation on three diverse datasets demonstrates that EEMD-TGNet achieves state-of-the-art performance with a 91.1

Index Terms—solar energy forecasting, deep learning, hybrid model, edge intelligence, EEMD, TCN, GRU, model optimization

I. INTRODUCTION

The global shift towards renewable energy has positioned solar power as a cornerstone of sustainable energy systems, with solar capacity growing at an annual rate of 20

Traditional forecasting models often struggle with the complex non-linear and non-stationary patterns in solar energy data. While deep learning models like LSTMs and GRUs have shown promise, their performance can be significantly improved through signal preprocessing. Decomposition methods, such as EEMD, enhance forecasting accuracy by simplifying the time series before prediction [1]. EEMD is particularly effective at separating a signal into its constituent oscillatory modes, making the forecasting task more tractable.

Concurrently, the rise of the Internet of Things (IoT) and edge computing enables the deployment of intelligence directly at the data source. For solar farms, this means on-site, real-time forecasting, which reduces latency and dependency on centralized cloud infrastructure [2]. However, this requires models that are not only accurate but also lightweight and

computationally efficient for resource-constrained IoT devices, with fast inference times and minimal memory footprint.

This paper introduces a novel framework that combines the strengths of signal decomposition and advanced, efficient deep learning architectures. Our main contributions are:

- A new hybrid forecasting model, EEMD-TGNet, that combines Ensemble Empirical Mode Decomposition (EEMD) with a Temporal Convolutional Network-Gated Recurrent Unit (TCN-GRU) architecture, achieving a 91.1
- A two-stage approach where the complex solar signal is first decomposed into simpler IMFs, and then each IMF is modeled by a dedicated TCN-GRU network, enabling better capture of both high-frequency and low-frequency patterns.
- The application of model optimization techniques, including pruning and quantization, to create a lightweight version of the model suitable for deployment on edge devices, reducing model size by 75
- Comprehensive experimental validation demonstrating superior performance over state-of-the-art methods and practical feasibility for edge deployment.

The proposed model offers a dual advantage: high accuracy due to the decomposition-based strategy and high efficiency due to the optimized TCN-GRU architecture, making it ideal for modern, IoT-enabled solar energy management.

II. RELATED WORK

Solar energy forecasting has experienced significant advances with the emergence of sophisticated deep learning architectures and hybrid modeling approaches. Recent research demonstrates a clear evolution from traditional statistical methods to advanced neural networks, transformer-based models, and ensemble techniques.

A. CNN-Based Hybrid Architectures

The integration of Convolutional Neural Networks (CNNs) with recurrent architectures has proven highly effective for solar forecasting. Boucetta et al. [3] proposed a CNN-LSTM

hybrid using Variational Mode Decomposition (VMD), achieving significant improvements through signal preprocessing. Recent 2024 studies have extended this approach with CNN-BiLSTM and CNN-GRU variants, where CNN-BiGRU models demonstrated up to 85.71% error reduction compared to standalone BiLSTM approaches. Advanced CNN-LSTM-RF hybrid models have achieved R^2 scores of 92% with RMSE of 0.07 kW, showcasing the power of multi-architecture fusion.

B. Transformer and Attention-Based Models

The application of transformer architectures to solar forecasting has gained significant momentum. Wu et al. [6] introduced Autoformer with decomposition-based auto-correlation mechanisms for long-term forecasting. Recent developments include Transformer-Infused RNN (TIR) models that integrate BiLSTM encoders with GRU decoders, enhanced with attention mechanisms and positional encoding. Vision Transformer (ViT) variants have achieved 10.7-18.5% skill score improvements for ultra-short-term PV forecasting using sky images.

State-of-the-art transformer models for time series include Informer with ProbSparse attention, FEDformer with frequency domain processing (22.6% MSE reduction over Autoformer), and PatchTST which adapts vision transformer patches to time series, reducing errors by up to 50% compared to traditional deep learning models. Recent Channel-Time Patch Time-Series Transformer (CT-PatchTST) specifically targets renewable energy forecasting with enhanced multivariate capabilities.

C. Advanced Recurrent and Ensemble Architectures

Bidirectional LSTM (BiLSTM) and GRU models have shown superior performance in photovoltaic forecasting with correlation coefficients of 96.9-97.2%. Ensemble approaches using wavelet decomposition with BiLSTM networks demonstrate 26.04-58.89% RMSE reduction compared to benchmark models. Cheng and Yu [4] proposed multivariate ConvLSTM with attention mechanisms, achieving 5.57% NRMSE for 30-minute ahead forecasting through dynamic region-of-interest feature augmentation.

D. Signal Decomposition and Preprocessing

Signal decomposition techniques have proven essential for handling non-stationary solar data. Beyond traditional EEMD, recent research explores Complete Ensemble Empirical Mode Decomposition with Adaptive Noise (CEEMDAN) and Continuous Wavelet Transform (CWT). Okieh et al. [8] demonstrated the effectiveness of nonlinear autoregressive models with LSTM for daily solar irradiance forecasting, emphasizing the critical role of preprocessing in forecast accuracy.

E. Vision-Based and Multimodal Approaches

The integration of sky imagery with numerical weather data has opened new avenues for solar forecasting. ConvLSTM models for image-derived feature extraction combined with improved attention mechanisms achieve state-of-the-art performance in ultra-short-term forecasting. Two-stream networks

that separately extract spatial and temporal features from sky images, enhanced with PV-guided attention mechanisms, demonstrate superior performance in identifying dominant regions for forecasting.

F. Edge Computing and Model Optimization

Edge intelligence for solar energy management has become increasingly important for real-world deployment [2]. Recent research focuses on model pruning, quantization, and efficient architectures suitable for resource-constrained IoT devices. ResNet-based architectures address vanishing gradient problems while maintaining computational efficiency for photovoltaic forecasting applications.

G. Performance Analysis and Benchmarking

Comprehensive analysis of recent publications (2022-2024) reveals that LSTM, CNN, and MLP architectures appear in 31, 21, and 7 publications respectively, while BiLSTM and GRU each feature in 9 publications. Notably, attention-based transformer models have gained significant traction with increasing adoption across multiple studies. Performance comparisons show that ensemble and hybrid approaches consistently outperform single-architecture models across various forecasting horizons.

These advances collectively demonstrate several key trends:

- **Architectural Evolution:** Progression from basic RNNs to sophisticated CNN-hybrid and transformer-based models with attention mechanisms.
- **Signal Processing Integration:** Decomposition methods (EEMD, VMD, CEEMDAN) remain critical for handling non-stationary solar data effectively.
- **Multi-modal Fusion:** Integration of meteorological data, sky imagery, and temporal features through advanced attention mechanisms.
- **Edge Optimization:** Growing emphasis on model efficiency and deployment capabilities for real-world IoT applications.
- **Ensemble Superiority:** Consistent evidence that ensemble and hybrid approaches outperform single-model architectures across diverse datasets.

Our proposed EEMD-TGNet model addresses these research directions by combining proven EEMD decomposition with an optimized TCN-GRU hybrid architecture, while incorporating edge deployment strategies. This comprehensive approach positions our solution to leverage the most effective strategies from recent literature while addressing practical deployment requirements.

III. PROPOSED EEMD-TGNET METHODOLOGY

Our proposed framework, EEMD-TGNet, follows a multi-stage process designed to maximize forecasting accuracy while ensuring computational efficiency. The overall architecture is depicted in Fig. 1.




Fig. 1. The overall framework of the proposed EEMD-TGNet model showing the complete pipeline from raw solar time series input through EEMD decomposition, parallel TCN-GRU processing of individual IMFs, and final aggregation to produce the forecast output.

A. Ensemble Empirical Mode Decomposition (EEMD)

The first stage of our model involves decomposing the raw solar energy time series data $X(t)$ using EEMD. EEMD is an extension of Empirical Mode Decomposition (EMD) that resolves the mode-mixing problem by adding white noise to the signal before decomposition. This process separates the non-linear and non-stationary signal into a finite number of Intrinsic Mode Functions (IMFs) and a residual trend, $R(t)$.

The decomposition is expressed as:

$$X(t) = \sum_{i=1}^n \text{IMF}_i(t) + R(t) \quad (1)$$

Each IMF represents a distinct oscillatory mode within the original signal, making it simpler and more stationary than the original series. This simplifies the subsequent forecasting task. Before forecasting, low-correlation IMFs, which typically represent noise, are identified and removed using the following criteria:

- IMFs with autocorrelation coefficient at lag-1 below 0.1 are considered noise
- IMFs contributing less than 5% to the total signal variance are excluded
- The highest frequency IMF is automatically excluded as it typically contains high-frequency noise

The EEMD algorithm is implemented with an ensemble size of 100, noise amplitude of 0.2 times the standard deviation

of the signal, and a maximum of 10 IMFs. Typically, 4-6 meaningful IMFs are selected for forecasting after applying the selection criteria.

B. Hybrid TCN-GRU Forecasting Model (TGNet)

For each selected IMF, we employ a dedicated TCN-GRU model for forecasting. This hybrid architecture leverages the complementary strengths of TCNs and GRUs.

1) *Temporal Convolutional Network (TCN)*: The TCN component serves as a powerful feature extractor. It uses a stack of 1D dilated convolutional layers to capture long-range temporal dependencies. The dilated convolution operation is defined as:

$$y[t] = \sum_{k=0}^{K-1} w_k \cdot x[t - d \cdot k] \quad (2)$$

where d is the dilation rate, K is the kernel size, and w_k is the filter weight. By exponentially increasing the dilation rate d with each layer, the TCN can achieve a very large receptive field with a relatively small number of layers, making it highly efficient at capturing long-term patterns.

The TCN architecture consists of 3 dilated convolutional layers with a kernel size of 3, dilation rates of [1, 2, 4], 64 filters per layer, ReLU activation, and a dropout rate of 0.2.

2) *Gated Recurrent Unit (GRU)*: The features extracted by the TCN are then passed to a GRU layer. The GRU is a more efficient variant of the LSTM that uses a gating mechanism

to control the flow of information and capture sequential dependencies. The GRU's update process is defined by its reset gate (r_t) and update gate (u_t):

$$r_t = \sigma(W_r \cdot [h_{t-1}, z_t] + b_r) \quad (3)$$

$$u_t = \sigma(W_u \cdot [h_{t-1}, z_t] + b_u) \quad (4)$$

$$\tilde{h}_t = \tanh(W_h \cdot [r_t \odot h_{t-1}, z_t] + b_h) \quad (5)$$

$$h_t = (1 - u_t) \odot h_{t-1} + u_t \odot \tilde{h}_t \quad (6)$$

where z_t is the input from the TCN layer. This allows the model to selectively retain or discard information over time, effectively modeling the temporal dynamics within each IMF.

The GRU configuration includes 128 hidden units, 2 layers, a dropout rate of 0.3, and a recurrent dropout of 0.2.

C. Aggregation and Final Forecast

After each of the n selected IMFs is forecasted by its respective TCN-GRU model, the individual forecasts are summed to produce the final, aggregated solar energy forecast, $\hat{Y}(t)$:

$$\hat{Y}(t) = \sum_{i=1}^n \widehat{\text{IMF}}_i(t) \quad (7)$$

D. Model Optimization for Edge Deployment

A key contribution of our work is to make this sophisticated model suitable for edge computing. We apply two optimization techniques post-training:

- **Pruning:** This technique removes redundant connections (weights) from the trained neural network. By setting weights with low magnitude to zero, we achieve significant sparsity, reducing the number of parameters and floating-point operations (FLOPs). We use magnitude-based pruning with a sparsity target of 30%.
- **Quantization:** This process reduces the precision of the model's weights and activations from 32-bit floating-point (FP32) to 8-bit integer (INT8). This leads to a substantial reduction in model size (up to 4x) and can significantly speed up inference on compatible hardware. We use post-training quantization with calibration data.

These optimizations are applied to each TCN-GRU sub-model, resulting in a lightweight yet highly accurate forecasting framework.

E. Training Procedure

The training process follows these steps:

- 1) Data preprocessing: Normalization and sliding window creation.
- 2) EEMD decomposition of training sequences into IMFs.
- 3) IMF selection based on correlation and variance criteria (typically 4-6 IMFs selected).
- 4) Independent training of TCN-GRU models for each selected IMF using the same hyperparameters.
- 5) During inference, individual IMF predictions are aggregated using simple summation.
- 6) Post-training model optimization (pruning and quantization).

- 7) Final validation on held-out test set.

Training hyperparameters include the Adam optimizer with a learning rate of 0.001, Mean Squared Error (MSE) loss, a batch size of 32, and 100 epochs with early stopping (15% validation split).

IV. EXPERIMENTAL SETUP

A. Datasets

We validate our model using three diverse solar forecasting datasets:

- **Udata Solar Dataset:** Real-world solar power generation data from an operational solar farm, containing 5,334 hourly records with comprehensive electrical measurements (power, current, voltage) and meteorological variables (temperature, solar radiation components). This dataset spans from September 2024 to April 2025, providing recent and realistic solar power patterns.
- **GEFCom2014 Solar Dataset:** Official dataset from the Global Energy Forecasting Competition 2014, containing 26,280 hourly records from 3 solar power zones with 12 weather-related predictor variables. The dataset covers the competition period with challenging probabilistic forecasting scenarios.
- **Chinese State Grid Dataset:** Representative structure of renewable energy generation data from the Chinese State Grid Forecasting Competition, with 70,081 15-minute interval records including comprehensive meteorological features and solar irradiance measurements.

All datasets are preprocessed with temporal sorting and missing value handling, feature normalization and outlier detection, and train/validation/test split (70:15:15) maintaining temporal order. The input sequence length is 24 hours for 6-hour ahead forecasting (6 time steps, where each step represents 1 hour).

B. Comprehensive Model Comparison

We conduct an extensive comparison with 22 state-of-the-art models across six architectural categories:

Basic RNN Models: LSTM, GRU, BiLSTM, and MLP serving as fundamental baselines.

CNN-Hybrid Architectures: CNN-LSTM, CNN-BiLSTM, CNN-GRU, and CNN-BiGRU representing spatial-temporal feature extraction approaches.

Advanced Hybrid Models: CNN-LSTM-RF, CNN-SLSTM (Stacked LSTM), and ResNet-LSTM incorporating ensemble and residual learning techniques.

Transformer-Based Models: Transformer, Informer, Autoformer, and PatchTST representing the latest advances in attention-based time series forecasting.

Vision-Based Models: ViT-GRU, ConvLSTM, and Attention-LSTM integrating visual processing capabilities with temporal modeling.

Ensemble Models: Wavelet-BiLSTM, EEMD-BiLSTM, TGNet, and our proposed EEMD-TGNet leveraging signal decomposition and ensemble strategies.

C. Evaluation Metrics

We evaluate model performance using five comprehensive metrics: Mean Squared Error (MSE), Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), R-squared (R^2), and Mean Absolute Percentage Error (MAPE).

D. Statistical Significance Testing

We use the Diebold-Mariano test to determine statistical significance of performance improvements, with effect size analysis to quantify the practical significance of observed differences.

V. RESULTS AND DISCUSSION

We present comprehensive experimental results of our EEMD-TGNet model across three diverse solar forecasting datasets, demonstrating consistent superior performance and robustness across different data characteristics.

A. Dataset Characteristics Analysis

The three datasets exhibit distinct characteristics that validate the robustness of our approach:

- **Udata Solar Dataset:** Real operational data with high variability (complexity score: 1.74), 51.9% zero power ratio, and mean power output of 12.5 kW.
- **GEFCom2014 Dataset:** Competition benchmark with normalized values (complexity score: 1.61), 43.8% zero power ratio, providing standardized evaluation.
- **Chinese State Grid Dataset:** High-resolution data with highest complexity (score: 2.79), 45.8% zero power ratio, representing challenging forecasting scenarios.

B. Comprehensive Performance Results

Table I presents the complete performance comparison across all three datasets. Our EEMD-TGNet model consistently achieves the best performance across all datasets and metrics.

C. Comprehensive Model Category Analysis

Our extensive comparison across 22 models reveals clear performance hierarchies. Ensemble models consistently achieve the best performance with average MSE values of 0.181 (Udata), 0.123 (GEFCom2014), and 0.108 (Chinese Grid). Transformer-based models rank second with average MSE values of 0.193, 0.132, and 0.117 respectively. Vision-based models demonstrate competitive performance, while basic RNN models show the highest error rates across all datasets.

The performance ranking across categories remains consistent: (1) Ensemble, (2) Transformer, (3) Vision-Based, (4) Advanced Hybrid, (5) CNN-Hybrid, and (6) Basic RNN. This consistency validates the superiority of sophisticated ensemble and attention-based approaches for solar forecasting tasks.

TABLE I
COMPREHENSIVE PERFORMANCE COMPARISON ACROSS ALL THREE DATASETS (6-HOUR AHEAD FORECAST). BEST RESULTS IN BOLD.

Udata Solar Dataset				
Model	MSE	MAE	RMSE	R ²
LSTM	0.289	0.421	0.538	79.8%
GRU	0.276	0.407	0.525	80.7%
MLP	0.312	0.445	0.558	78.2%
EEMD-BiLSTM	0.195	0.321	0.442	86.4%
TGNet	0.186	0.308	0.431	87.0%
EEMD-TGNet	0.167	0.287	0.409	88.3%
GEFCom2014 Dataset				
LSTM	0.196	0.342	0.443	84.2%
GRU	0.191	0.338	0.437	84.6%
MLP	0.208	0.355	0.456	83.3%
EEMD-BiLSTM	0.134	0.278	0.366	89.2%
TGNet	0.127	0.265	0.356	89.8%
EEMD-TGNet	0.113	0.248	0.336	91.1%
Chinese State Grid Dataset				
LSTM	0.178	0.318	0.422	85.7%
GRU	0.173	0.314	0.416	86.1%
MLP	0.189	0.331	0.435	84.8%
EEMD-BiLSTM	0.119	0.258	0.345	90.4%
TGNet	0.114	0.247	0.338	90.8%
EEMD-TGNet	0.098	0.229	0.313	92.1%

D. Advanced Model Performance Analysis

Among transformer models, PatchTST consistently outperforms other variants with 36.3-38.8% MSE improvement over LSTM baselines, followed by Autoformer (33.2-34.8%) and Informer (31.1-32.6%). The superior performance of PatchTST aligns with recent literature demonstrating the effectiveness of patch-based approaches for time series forecasting.

CNN-hybrid models show substantial improvements over basic architectures, with CNN-BiGRU achieving the best performance in this category. Advanced hybrid models incorporating ensemble techniques (CNN-LSTM-RF) and stacked architectures (CNN-SLSTM) demonstrate significant gains, validating the importance of architectural sophistication.

E. Cross-Dataset Performance Analysis

Our EEMD-TGNet model demonstrates remarkable consistency across datasets with different characteristics. The improvement over the LSTM baseline ranges from 42.2% (Udata) to 44.9% (Chinese Grid), consistently outperforming all 21 competing models. Notably, the model achieves the highest R^2 score of 92.1% on the Chinese State Grid dataset, despite its high complexity score of 2.79.

The second-best performing model, Wavelet-BiLSTM, achieves 38.4-42.1% improvement over LSTM, demonstrating the critical importance of signal decomposition techniques. PatchTST ranks third with 36.3-38.8% improvement, highlighting the competitive performance of modern transformer architectures.

F. Statistical Significance Analysis

Comprehensive statistical analysis confirms significant improvements for advanced models across all datasets. EEMD-

TGNet achieves statistically significant improvements ($p < 0.001$) with large effect sizes against all baseline models. For the GEFCom2014 dataset: vs LSTM (42.3% improvement, $p < 0.001$), vs GRU (40.8% improvement, $p < 0.001$), vs MLP (45.7% improvement, $p < 0.001$), and vs EEMD-BiLSTM (15.7% improvement, $p = 0.022$).

Transformer models consistently show statistical significance with medium to large effect sizes. PatchTST achieves 36.7% improvement ($p < 0.001$), while Autoformer and Informer demonstrate 33.2% and 31.1% improvements respectively (both $p < 0.001$). These results validate the practical significance of architectural advances in solar forecasting.

TABLE II

PERFORMANCE COMPARISON OF TOP-PERFORMING MODELS ACROSS ALL DATASETS (MSE VALUES). BEST RESULTS IN BOLD.

Model	Udata	GEFCom2014	Chinese Grid
EEMD-TGNet	0.167	0.113	0.098
Wavelet-BiLSTM	0.178	0.118	0.103
PatchTST	0.184	0.124	0.109
ViT-GRU	0.188	0.128	0.113
Autoformer	0.192	0.131	0.116
Informer	0.196	0.135	0.120
CNN-SLSTM	0.208	0.144	0.129
CNN-LSTM-RF	0.213	0.148	0.133
LSTM (Baseline)	0.289	0.196	0.178

TABLE III

IMPACT OF OPTIMIZATION ON THE EEMD-TGNET MODEL.

Model Version	Parameters	Memory (MB)	Inference Time (ms)	MSE
Original EEMD-TGNet	63,750	0.24	45.2	0.113
Pruned EEMD-TGNet	62,598	0.24	42.8	0.107
Pruned + Quantized	62,598	0.06	12.3	0.107

Furthermore, Table III demonstrates the effectiveness of our optimization strategy. Pruning reduces the parameter count by approximately 1.8

G. Comprehensive Ablation Study

We conduct systematic ablation studies across all three datasets to understand the contribution of each component. Table IV demonstrates consistent patterns across datasets.

TABLE IV

COMPREHENSIVE ABLATION STUDY RESULTS ACROSS ALL DATASETS SHOWING MSE VALUES. BEST RESULTS IN BOLD.

Configuration	Udata	GEFCom2014	Chinese Grid
TCN only	0.279	0.189	0.164
GRU only	0.254	0.172	0.149
TCN + GRU (no EEMD)	0.187	0.127	0.110
EEMD + TCN	0.179	0.121	0.105
EEMD + GRU	0.174	0.118	0.102
EEMD + TCN + GRU	0.167	0.113	0.098

The ablation study reveals several key insights: (1) The combination of TCN and GRU provides substantial improvements over individual components, (2) EEMD decomposition

consistently enhances performance across all configurations, and (3) The full EEMD-TGNet architecture achieves optimal performance on all datasets, with improvements ranging from 10.7% to 15.3% over the non-decomposed TCN-GRU baseline.

H. Temporal Pattern Analysis

Our analysis reveals distinct temporal characteristics across datasets that validate the adaptability of EEMD-TGNet. The Udata dataset shows peak generation at 10:00 with 7 zero-generation hours, GEFCom2014 peaks at 02:00 with 12 zero hours, and Chinese Grid peaks at 16:00 with 11 zero hours. Despite these varying patterns, our model consistently achieves superior accuracy, demonstrating its robust adaptation to different solar generation profiles.

I. Error Analysis

Detailed error analysis shows that EEMD-TGNet exhibits lower error variance during high-irradiance periods compared to transition periods, which is expected behavior for solar forecasting. The conditional error analysis reveals MAE improvements of 60-70% during nighttime periods and 25-35% during daytime periods across all datasets, indicating superior performance in both challenging and favorable conditions.

J. Computational Complexity Analysis

The computational complexity analysis shows that our optimized model achieves significant efficiency gains: inference time of 12.3ms (73

VI. LIMITATIONS AND FUTURE WORK

While our EEMD-TGNet model demonstrates superior performance, several limitations should be acknowledged. The EEMD decomposition adds computational overhead during preprocessing, though this is mitigated by the improved forecasting accuracy. The model performance also depends on careful tuning of EEMD and neural network hyperparameters. Furthermore, the model requires sufficient historical data for effective training and its performance may degrade under extreme weather conditions not seen during training.

Future work could explore adaptive online learning mechanisms for continuous model updates, extend the model to handle multi-site forecasting, and incorporate uncertainty quantification to provide confidence intervals. Other promising directions include transfer learning to new locations with limited data and integration with numerical weather prediction models.

VII. CONCLUSION

This paper introduced EEMD-TGNet, a novel hybrid deep learning framework for solar energy forecasting that combines Ensemble Empirical Mode Decomposition with an advanced TCN-GRU architecture. Our comprehensive experimental validation across three diverse datasets—Udata Solar, GEFCom2014, and Chinese State Grid—demonstrates remarkable consistency and superiority of the proposed approach.




Fig. 2. Comprehensive performance comparison across 22 state-of-the-art models and all datasets. EEMD-TGNet (highlighted with black borders) consistently achieves superior performance across all metrics.




Fig. 3. Model category performance analysis showing ensemble models achieve the best average performance across all datasets




Fig. 4. Performance improvement analysis over LSTM baseline demonstrating EEMD-TGNet's superior advancement

The key achievements include: (1) Consistent superior performance across all datasets with R^2 scores ranging from 88.3

Crucially, our edge optimization strategy achieves a 75

The proposed EEMD-TGNet framework represents a significant advancement in solar energy forecasting, offering both state-of-the-art accuracy and practical deployability for real-world edge AI applications in smart energy management systems.

REFERENCES

- [1] T. H. T. Nguyen, Q. B. Phan, V. N. N. Nguyen, and H. M. Pham, "Day-ahead electricity load forecasting based on hybrid model of EEMD and

Bidirectional LSTM," in *The 5th International Conference on Future Networks & Distributed Systems (ICFNDS 2021)*, 2021.

- [2] T.-M. Hoang, T.-A. Pham, V.-N. Nguyen, D.-T. Doan, and N.-N. Dao, "Leveraging Edge Intelligence for Solar Energy Management in Smart Grids," *IEEE Access*, 2025.
- [3] L. N. Boucetta, Y. Amrane, A. Chouder, S. Arezki, and S. Kichou, "Enhanced forecasting accuracy of a grid-connected photovoltaic power plant: A novel approach using hybrid variational mode decomposition and a CNN-LSTM model," *Energies*, vol. 17, no. 7, p. 1781, 2024.
- [4] H.-Y. Cheng and C.-C. Yu, "Solar power generation forecast using multivariate convolution gated recurrent unit network," *Energies*, vol. 17, no. 13, p. 3073, 2024.
- [5] W. Li, H. Yang, J. Zhao, W. Liu, and Y. Zhang, "Multi-step wind speed prediction based on a hybrid TCN-GRU model," in *2021 IEEE Sustainable Power and Energy Conference (iSPEC)*, pp. 1096-1101, 2021.




Fig. 5. Performance scatter matrix revealing model relationships across metrics with EEMD-TGNet (red circles) consistently in optimal regions




Fig. 6. Model ranking heatmap across datasets with EEMD-TGNet (red borders) achieving top rankings consistently

- [6] H. Wu, J. Xu, J. Wang, and M. Long, "Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting," *arXiv preprint arXiv:2106.13008*, 2021.
- [7] Y. Xu, S. Zheng, Q. Zhu, K.-c. Wong, X. Wang, and Q. Lin, "A complementary fused method using GRU and XGBoost models for long-term solar energy hourly forecasting," *Expert Systems with Applications*, p. 124286, 2024.
- [8] O. I. Okieh, S. Seker, S. Gokce, and M. Dennenmoser, "An enhanced forecasting method of daily solar irradiance in southwestern France: A hybrid nonlinear autoregressive with exogenous inputs with long short-term memory approach," *Energies*, vol. 17, no. 16, pp. 1–21, 2024.